

Clause IQ: Intelligent Legal Documents Processing Using Retrieval Augmented Generation

*A Project Report submitted in partial fulfilment of the requirements
for the award of the degree of*

Master of Computer Applications

by

Ayush Pathak	12584200052
Shreyansh Agarwal	12584200182
Ansh Pal Singh	12584200034
Rudransh Ramani	12584200164

Group No. II

Under the Guidance of

Dr. Sachendra Singh Chauhan

Department of Computer Engineering & Applications

Institute of Engineering & Technology



GLA University

Mathura- 281406, INDIA

April, 2026



Department of Computer Engineering and Applications
GLA University, 17 km Stone, NH#2, Mathura-Delhi Road,
P.O. Chaumuhan, Mathura-281406 (U.P.)

Declaration

I hereby declare that the work which is being presented in the M.C.A. Project **“ClauseIQ: Intelligent Legal Documents Processing Using Retrieval Augmented Generation”**, in partial fulfillment of the requirements for the award of the *Master of Computer Applications* and submitted to the Department of Computer Engineering and Applications of GLA University, Mathura, is an authentic record of our own work carried under the supervision of **Dr. Sachendra Singh Chauhan, Assistant Professor.**

The contents of this project report, in full or in parts, have not been submitted to any other Institute or University for the award of any degree.

Sign _____
Name: Shreyansh Agarwal
University Roll No.: 12584200182

Sign _____
Name: Ansh Pal Singh
University Roll No.: 12584200034

Sign _____
Name: Ayush Pathak
University Roll No.: 12584200052

Sign _____
Name: Rudransh Ramnani
University Roll No.: 12584200164

Certificate

This is to certify that the above statements made by the candidate are correct to the best of my/our knowledge and belief.

Supervisor
(Dr. Sachendra Singh Chauhan)
Designation of Supervisor
Dept. of Computer Engg, & App.

Project Co-ordinator
(Dr. Sachendra Singh Chauhan)
Assistant Professor
Dept. of Computer Engg, & App.

Program Co-ordinator
(Mr. Siddharth Singh)
Assistant Professor
Dept. of Computer Engg, & App.

Date:

ACKNOWLEDGEMENT

We would like to express our sincere gratitude to our supervisor, Dr. Sachendra Singh Chauhan, for his continuous guidance, valuable suggestions, and patient encouragement throughout the development of this project. The successful completion of this project has been possible due to his constant support, insightful feedback, and motivation at every stage of the work.

We also wish to acknowledge the faculty members of the Department of Computer Engineering & Applications, GLA University, for their academic guidance and for providing an environment that encouraged systematic learning and disciplined work.

We extend our heartfelt thanks to our families and friends for their understanding, encouragement, and unwavering support during the preparation of this report and the development of the application.

Group Members:

Shreyansh Agarwal
University Roll No.: 12584200182

Ansh Pal Singh
University Roll No.: 12584200034

Ayush Pathak
University Roll No.: 12584200052

Rudransh Ramnani
University Roll No.: 12584200164

ABSTRACT

Access to legal information remains a significant challenge for students, general users, and early-stage practitioners who frequently encounter difficulty interpreting legal language or locating relevant statutory provisions without professional assistance. Existing tools either require domain expertise to operate effectively or rely on cloud-based infrastructure that raises concerns about document privacy and data confidentiality. This project addresses these limitations through the design and development of ClauseIQ, an AI-assisted legal query and document analysis system built on a fully local processing pipeline.

The system employs Retrieval-Augmented Generation (RAG) as its core technique, wherein user queries expressed in natural language are matched against an indexed corpus of legal text using semantic similarity search powered by FAISS and the sentence-transformers library. The most relevant passages are retrieved and passed as context to a locally hosted language model via Ollama, which generates accurate, traceable, and context-aware responses without transmitting any data to external servers.

Beyond the core question-answering capability, the system incorporates user authentication with session management, PDF and DOCX document upload with automatic chunking and indexing, persistent conversation history, document-type detection, and a responsive web interface. The backend is implemented using the Flask framework with SQLite for structured data storage, maintaining a clean separation between user interaction, retrieval, and generation layers.

Experimental evaluation demonstrates that the RAG-based retrieval approach consistently outperforms keyword-matching baselines in response relevance across varied legal query types. The final system functions as a complete, deployable web application rather than a partial prototype. It is not intended to replace professional legal counsel but serves as a guided assistant that makes legal material significantly more accessible, explorable, and reviewable for non-specialist users.

Contents

Declaration	iii
Certificate	iv
Acknowledgement	v
Abstract	vi
List of Figures	vii
List of Tables	viii
List of Symbols, Abbreviations and Nomenclature	ix
Chapter 1 Introduction	1
1.1 Motivation and Overview	1
1.2 Problem Statement	2
1.3 Objectives of the Project	2
1.4 Summary of Similar Applications	3
1.5 Scope and Significance	4
1.6 Organization of the Project Report	5
Chapter 2 Software Requirement Analysis	5-11
Chapter 3 Software Design	12-21
Chapter 4 Implementation and User Interface	22-30
Chapter 5 Software Testing and Result Analysis	31-36
Chapter 6 Conclusion and Future Scope	37-42
Appendices	43-49
References	50

List of Figures

Figure 3.1 ClauseIQ System Architecture	20
Figure 3.2 Authentication and Session Flow	22
Figure 3.3 Document Ingestion Pipeline	23
Figure 4.1 Representative User Interface Layout	31
Figure 4.2 Landing Page	29
Figure 4.3 Legal AI Interface	29
Figure 4.4 Representative UI Layout	30
Figure 4.5 Geo-tagged Image of Project Team	30

List of Tables

Table 2.1 Functional Requirements	12
Table 2.2 Non-Functional Requirements	13
Table 2.3 Feasibility Analysis	14
Table 3.1 Major Modules and Responsibilities	19
Table 3.2 Database Schema Summary	24
Table 5.1 Black Box Test Cases	40
Table 5.2 Security and Robustness Test Cases	42
Table 6.1 Future Enhancement	43
Table 6.2 User / Session	44
Table 6.3 Document Metadata	45
Table 6.4 Document Chunks	46
Table 6.5 Embeddings	46
Table 6.6 Query Processing	47
Table 6.7 Retrieval Results	47
Table 6.8 Response Generation	48
Table 6.9 Error Handling	48

Chapter 1

Introduction

1.1 Motivation And Overview

Access to legal information has become increasingly important. Individuals, students, and early-stage practitioners often need to interpret statutes, understand procedural requirements, or review documents without immediate assistance from a legal professional. Even though legal text is available in digital form, the gap between raw statutory content and practical understanding remains large. Most users find legal material in dense, jargon-heavy documents that are hard to navigate without prior training.

ClauseIQ aims to close this gap by offering a conversational interface over indexed legal content. Instead of needing users to browse statutes manually or build precise search queries, the system lets them ask questions in plain language and receive answers based on relevant source material. This method is especially useful in the Indian legal context, where multiple overlapping codes—including the IPC, CrPC, CPC, and the recently enacted BNS—make it hard for non-specialists to quickly find the applicable provision.

The project also tackles a secondary issue: privacy. Most cloud-based legal AI tools process user queries and documents on remote servers, which may not be suitable for sensitive content. ClauseIQ processes all data locally. This ensures that documents uploaded by users and queries entered into the system never leave their own machine.

From an academic point of view, the project shows that retrieval-augmented generation, vector search, and local language model integration can be combined into a coherent and usable system as part of an MCA final year project.

1.2 Problem Statement

The problem ClauseIQ addresses is not the lack of legal text-statutes, codes, and procedural rules are widely available in digital form. The real issue is the lack of an effective way to interact with that text. A user who doesn't know which section of the CrPC covers bail applications, or what the BNS replaced, cannot benefit from access to the full text unless they already know where to look.

Traditional search tools return documents instead of reasoned answers. A keyword search for "bail conditions" might return dozens of sections from various acts, requiring the user to read and interpret each one separately. This process is time-consuming, error-prone, and discouraging for those without legal training.

The problem is even worse when users have their own documents-property deeds, rental agreements, legal notices-that they need to understand. No general-purpose search tool can answer questions about a specific uploaded document. The user must either read the document themselves or ask a lawyer for even basic clarification.

ClauseIQ addresses both facets of this problem: general legal queries over a statutory corpus and document-specific questions about user-uploaded files. The system identifies the most relevant passages from either source and presents them through a conversational interface, making it easier to go from question to answer.

1.3 Objectives Of The Project

The project is guided by the following objectives:

1. **Semantic Retrieval:** To implement a retrieval pipeline that matches user queries against indexed legal text using vector similarity rather than keyword overlap. This improves the relevance of retrieved passages.
2. **Context-Aware Answer Generation:** To integrate a locally hosted language model that generates answers based on the retrieved context. This ensures that responses are tied to source material rather than generated from the model's memory.
3. **Document Upload and Indexing:** To support user-uploaded PDF and DOCX files so that the system can also answer questions about specific documents in addition to the general legal corpus.
4. **User Authentication and Session Management:** To implement a secure login, ensure per-user document isolation, and create persistent conversation history. This allows

the system to function as a complete multi-user application rather than a single-session prototype.

5. **Responsible Design:** To build the system transparently regarding its limitations. It should clearly distinguish between AI-generated assistance and professional legal advice, and all processing should remain local to protect user privacy.

1.4 Summary Of Similar Applications

Several existing tools provide legal search or document assistance, but each has significant limitations that ClauseIQ aims to address.

Harvey AI and Spellbook are among the leading legal AI platforms currently available. They offer contract review, clause drafting, and document analysis with high accuracy. However, both are designed exclusively for legal professionals, priced for law firms, and process all data on remote cloud infrastructure. Neither is suitable for students, general users, or those who prioritize privacy.

DoNotPay offers a consumer-facing legal assistant interface with guided workflows for common legal tasks. Its coverage is limited to specific use cases and primarily focuses on US and UK law, with little attention to Indian statutes or procedures.

Indian Kanoon is a popular legal search engine in India that indexes court judgments and statutes. It is effective for keyword-based case law research but returns raw documents instead of synthesized answers, and it does not support conversational querying or document uploads.

ChatGPT and similar general-purpose LLMs can answer legal questions in a conversational style but are not based on a verified corpus. Their answers might sound plausible, but they cannot be traced back to source material. Additionally, there is a risk of uploading the sensitive documents to the cloud/online storage.

ClauseIQ fills a unique role: it integrates conversational querying, document-specific analysis, and fully local processing in one deployable application. No other reviewed tools offer this combination.

1.5 Scope And Significance

ClauseIQ covers several functional areas: natural language querying over an indexed legal corpus, answering questions about user-uploaded documents, secure user authentication with data isolation, persistent conversation history, and a responsive web interface suitable for regular use.

The system intentionally avoids activities that require professional judgment-it does not draft legal documents, predict case outcomes, or provide advice on specific disputes. This boundary is clearly communicated to users within the interface and reinforced by the design, which always ties responses back to source passages instead of making unsupported conclusions.

The project is significant for two reasons. Practically, it provides a working tool that makes Indian legal material more accessible to a non-specialist audience, including students studying law-adjacent subjects, individuals dealing with property or tenancy issues, and early-stage practitioners needing a quick reference tool. Technically, it demonstrates that a retrieval-augmented generation-based legal assistant can be built and used without cloud dependency, proprietary APIs, and within the limits of a typical development machine.

The project also lays a foundation for future expansions, such as citation linking to official government legal databases, role-based access for institutional deployment, analytics on query patterns, and multilingual support for Hindi-language queries.

1.6 Organization Of The Project Report

Chapter 2: Software Requirement Analysis identifies the functional and non-functional requirements of the system, assesses technical and operational feasibility, and defines the modules and use cases that guide the design.

Chapter 3: Software Design presents the system architecture, data flow diagrams, UML models including class and sequence diagrams, and the database schema. This chapter explains how the retrieval, generation, and storage components are structured and how they interact.

Chapter 4: Implementation and User Interface describes the development environment, the technology stack, and the implementation of each module. It includes screenshots of the working application and clarifies key implementation choices.

Chapter 5: Software Testing documents the test cases applied to the system using both black-box and white-box testing strategies, along with the results of each test.

Chapter 6: Conclusion summarizes the outcomes of the project, reflects on how well each objective was met, and highlights directions for future development.

Chapter 2

Software Requirement Analysis

2.1 Stakeholders And User Roles

ClauseIQ is designed to serve more than one category of user, and the distinct needs of each category directly influence the system design. Three primary stakeholder groups have been identified, each with different expectations from the application.

General users form the largest group. These are individuals with no formal legal background - students, property owners, tenants, or anyone seeking to understand a legal document or provision. This group requires a simple interface, plain-language responses, and minimal technical setup. They interact with the system primarily through the chat interface and the document upload facility, and their primary concern is whether they can obtain a useful answer quickly without needing to understand how the system operates internally.

Legal students and early practitioners represent a second group with somewhat greater familiarity with legal terminology. They use the system for study, reference, or preliminary research, and they benefit particularly from accurate retrieval, traceable source passages, and the ability to query specific documents they are already working with. For this group, the quality and grounding of the response carries greater importance than interface simplicity.

2.2 Functional Requirements

The functional requirements of ClauseIQ define the visible capabilities that the system must provide across each of its modules. Together they describe the complete set of user-facing and system-level behaviours that the application is expected to support.

The chat and query module carries the primary intelligence of the application. It must accept natural-language queries from authenticated users, retrieve the most semantically relevant passages from the indexed corpus, pass those passages as context to the locally hosted language model, and return a coherent, source-grounded response.

Before processing any query, the system must screen for prompt injection or jailbreak patterns and classify the query as document-specific, general legal, or out-of-scope, routing each category through an appropriate response path.

Document management requires the system to accept PDF and DOCX uploads from authenticated users, extract readable text from each file, divide that text into overlapping chunks of defined size, generate vector embeddings for each chunk, and add those embeddings to the persistent FAISS index. The system must track the indexing status of each document and communicate it to the user, and it must allow users to delete any uploaded document along with all associated chunks, embeddings, and metadata so that no residual data persists after deletion.

Conversation history management requires that every query and its corresponding response be stored against the authenticated user's account and organised into named sessions. Past sessions must be retrievable through a sidebar, named automatically from the initial query, and resumable by the user at any subsequent login without loss of continuity.

2.3 Non-Functional Requirements

Non-functional requirements govern the quality attributes of the system - the properties that determine how well it performs its functions rather than what those functions are. For ClauseIQ, the most significant non-functional dimensions are performance, usability, and maintainability.

Performance is a concern primarily in the retrieval step. As the document index grows with additional uploads and seeded legal data, the time required to identify relevant passages must remain within an acceptable response window. The FAISS in-memory index is selected specifically because it supports fast similarity search across large numbers of vectors without requiring a dedicated external service. The background threading model used during document indexing ensures that upload processing does not block the user interface while indexing is in progress.

Security requirements cover authentication, data access, and session management. User passwords must be stored as cryptographically hashed values and must never be persisted in plain text. Sessions must be managed through server-side tokens with

defined expiry periods so that abandoned sessions do not persist indefinitely. Per-user document isolation must be enforced at the database level so that a retrieval query from one user cannot surface documents or conversation history belonging to another. Failed login attempts must be tracked and the account temporarily locked after a configurable threshold is exceeded, reducing exposure to automated credential attacks.

Privacy is treated as a primary requirement rather than a secondary consideration. Because the system is designed for local deployment, no user query, uploaded document, or conversation record is transmitted to any external server at any point in the processing pipeline. The language model runs through Ollama on the user's own machine, and all persistent data is stored in a local SQLite database. This makes the system appropriate for users who handle sensitive personal or professional documents and are unwilling to accept the data exposure associated with cloud-based alternatives.

Usability requirements specify that the interface must be navigable without prior technical training, that error messages must be informative and expressed in plain language rather than system codes, and that the status of background operations such as document indexing must be visible to the user at all times. The system must also function correctly on standard consumer hardware without requiring a graphics processing unit or any specialised computing infrastructure, since the intended user base is unlikely to have access to high-performance resources.

Maintainability is addressed through the modular separation of the codebase into distinct service layers for authentication, document ingestion, vector retrieval, response generation, and conversation history. Each module exposes a defined interface to the others, which means that a component can be updated or replaced without requiring changes throughout the entire system. This is particularly relevant for the language model component, where the underlying model may be substituted as more capable options become available.

2.4 Feasibility Analysis

The feasibility study evaluates whether the project can be implemented successfully using the available tools, skills, and time frame. ClauseIQ is assessed across three dimensions: technical feasibility, operational feasibility, and economic feasibility.

From a technical perspective, the project is feasible because every component of the system is built on mature, well-documented open-source software. Flask provides the web framework, SQLAlchemy manages database interactions, FAISS handles vector indexing and retrieval, sentence-transformers generates embeddings, and Ollama provides local language model inference. None of these dependencies require paid licences, specialised hardware, or access to external APIs. Each subsystem has been validated independently before integration, and the combination has been shown to function correctly within the constraints of a standard development environment.

From an operational perspective, the system is feasible because the user-facing workflow is straightforward and requires no technical knowledge to execute. A user registers an account, logs in, uploads a document or enters a query, and receives a response. The complexity of the retrieval and generation pipeline is entirely encapsulated within the service layer and is not exposed to the user in any form. Administrative tasks are similarly confined to a dedicated dashboard that can be operated without knowledge of the underlying architecture.

From an economic perspective, the project is feasible because it avoids any dependency on paid inference services or proprietary APIs. The entire system runs on a standard development machine, and the only ongoing resource requirement is the local compute time needed to run the language model through Ollama. This makes the project accessible to academic institutions, individual researchers, and small organisations that cannot justify the recurring cost of cloud-based AI services, and it ensures that the system remains deployable without ongoing expenditure after initial setup.

2.5 Constraints And Assumptions

Every software project operates within a set of constraints that shape both the design and the expected behaviour of the final system. Acknowledging these constraints clearly is important because it establishes honest expectations for evaluators, users, and future developers.

The most significant technical constraint is the quality of the FAISS index, which depends directly on the quality of the text extraction and chunking pipeline. If an uploaded document consists of scanned images rather than machine-readable text, The

system communicates this limitation to the user through the document status indicator, which reflects indexing failure when extraction produces no usable text.

A second constraint concerns the response quality of the locally hosted language model. Smaller models that run on consumer hardware without a GPU produce responses that are noticeably less detailed and occasionally less precise than responses from large cloud-hosted models. The system mitigates this by ensuring that the retrieved context is always passed explicitly to the model, reducing reliance on the model's parametric knowledge and improving the grounding of generated responses. Nevertheless, users should be aware that output quality is bounded by the capability of the model running on their machine.

The system operates under several assumptions. It assumes that uploaded documents are legible, relevant to the legal domain, and provided in a supported file format. It also assumes that the Ollama service is running locally before the application is started, since the system does not attempt to launch or manage Ollama as a subprocess. A health check is provided within the interface to indicate whether the language model is available, but the system cannot function without it.

From a scope perspective, the system does not attempt to provide legal advice, predict case outcomes, or draft legal documents on behalf of users. These activities require professional judgment that no automated system can reliably replicate, and they fall outside the defined boundaries of this project.

2.6 Module Decomposition

The project is divided into six principal modules, each responsible for a clearly defined subset of the system's functionality. This separation reflects the principle of single responsibility, where each module can be developed, tested, and maintained without requiring detailed knowledge of the others.

The document ingestion module is responsible for receiving uploaded files, extracting text through format-appropriate parsers, splitting the extracted text into overlapping chunks of defined size, generating embeddings for each chunk using the sentence-transformers model, and persisting both the chunks and their embeddings to the database and the FAISS index. It also manages the document status lifecycle from

upload through indexing to completion or failure, and it handles deduplication by computing a hash of each file at upload time.

The vector retrieval module accepts a query string, generates a query embedding using the same model used during ingestion, and performs a similarity search against the FAISS index to identify the most semantically relevant chunks. It returns ranked results along with their source metadata, which are subsequently passed to the generation module.

The generation module constructs a structured prompt from the retrieved context and the user's original query, submits it to the locally hosted language model through the Ollama API, and returns the generated response. It enforces a guardrail check against prompt injection before passing the query forward and handles fallback behaviour gracefully when the model service is unavailable.

The admin module provides the administrator interface for monitoring system metrics, managing user accounts, and reviewing access requests. It enforces role-based access control so that its endpoints remain inaccessible to regular users regardless of how requests are constructed.

2.7 Use Case Scenarios

Use case scenarios translate the abstract requirements described in earlier sections into concrete descriptions of how users interact with the system under realistic conditions. Three representative scenarios are described below.

In the most common scenario, a general user visits the application for the first time, registers an account, and is directed to the chat interface. They submit a query such as "What is the punishment for cheating under the BNS?" The system screens the query for unsafe patterns, classifies it as a general legal query, retrieves the most relevant passages from the indexed BNS corpus, and passes them as context to the language model. The model returns a response identifying the applicable section and summarising its provisions. The user reads the response and closes the application. On a subsequent visit they log in again and find their previous conversation preserved and accessible from the sidebar.

In a document-specific scenario, a user uploads a rental agreement in PDF format. The system extracts the text, generates embeddings for each chunk, and notifies the user once indexing is complete. The user then submits the query "What is the notice period stated in this agreement?" The system retrieves the relevant clause from the indexed document, passes it to the language model, and returns a response that quotes the specific provision. The user may continue asking follow-up questions about the same document within the same session, and each subsequent response remains anchored to the uploaded content.

2.8 Data Requirements

The application depends on two categories of data: structured relational data managed through SQLite and vector data managed through the FAISS index.

The relational data layer stores user accounts, document metadata, document chunks, chunk embeddings, conversation sessions, and chat messages. Each entity carries a well-defined set of attributes. A document record stores the original filename, storage path, SHA-256 hash for deduplication, MIME type, file size, upload timestamp, and current indexing status. A chunk record stores the extracted text, its positional index within the parent document, and a metadata object containing the source filename and document identifier. A conversation session record stores the session name, creation timestamp, and a reference to the owning user account.

The vector data layer stores the numerical embeddings generated from each chunk. Embeddings are persisted both in the SQLite embeddings table as JSON arrays and in a binary FAISS index file on disk. The FAISS index is loaded into memory when the application starts and updated incrementally as new documents are indexed. Persisting the index to disk is essential because an in-memory-only index would be destroyed on every server restart, requiring all documents to be re-indexed before retrieval could resume.

Metadata integrity is as important as text quality for the correct functioning of the system. Source attribution in generated responses depends on chunk metadata being accurate and consistently structured. Document deduplication depends on the SHA-256 hash being computed at upload time before any processing begins.

Chapter 3

Software Design


3.1 Overall System Architecture

ClauseIQ follows a layered architecture in which the presentation layer, application layer, retrieval layer, and data layer remain logically separated. This separation ensures that each stage of the system can be understood, tested, and modified independently without introducing unintended effects in adjacent components.

The presentation layer consists of the web interface rendered in the user's browser. It communicates with the application layer exclusively through HTTP requests to a defined set of REST API endpoints. The application layer, implemented using the Flask framework, receives these requests, enforces authentication, coordinates the appropriate service modules, and returns structured responses. The retrieval layer handles the embedding of queries, the similarity search against the FAISS index, and the selection of relevant passages. The data layer encompasses the SQLite relational database, the FAISS binary index file, and the local file system used for document storage.

No layer communicates directly with a layer that is not immediately adjacent to it. The presentation layer has no access to the database, and the retrieval layer has no knowledge of the user interface. This constraint keeps the codebase comprehensible and creates clear boundaries within which future modifications can be made. For example, the retrieval layer could be replaced with a different vector search backend without requiring any changes to the presentation layer or the authentication module.

The architecture also reflects the privacy requirement established in Chapter 2. Because the language model is invoked through a local Ollama instance rather than a remote API, all data remains within the deployment environment at every stage of the pipeline.



RAG ARCHITECTURE - CLAUSEIQ

3.2 Document Ingestion Pipeline

The document ingestion pipeline prepares external material for semantic search and subsequent reuse during query processing. It is designed as a sequential, fault-tolerant process that moves a document from raw upload through text extraction, chunking, embedding generation, and index registration.

When a file is uploaded, the system first validates its format, verifies that it does not exceed the configured size limit, and computes a SHA-256 hash of its contents. If a document with the same hash already exists in the database, the upload is rejected as a duplicate and the user is informed accordingly. This deduplication step prevents redundant index entries and keeps the retrieval corpus free of repeated content.

If the document is accepted, it is saved to the local file system under a path that incorporates the hash value to ensure uniqueness. A document record is created in the database with a status of awaiting index, and the indexing process is dispatched to a background thread so that the upload response is returned to the user immediately without waiting for indexing to complete.

The background thread extracts text from the document using format-specific parsers. PDF files are processed using `pypdf`, which extracts the text content of each page up to a defined maximum. DOCX files are processed using `python-docx`, which iterates over the paragraph structure of the document. If text extraction yields no usable content, the system attempts OCR-based extraction using `pytesseract` as a fallback. If this also fails,

a placeholder chunk is created and the document status is updated to reflect that full-text analysis is not available.

When readable text is successfully extracted, it is divided into overlapping chunks of approximately eight hundred characters, with boundaries aligned to paragraph breaks where possible. Overlapping chunks are used rather than disjoint ones so that retrieval is not disrupted when a relevant passage spans a natural boundary. Each chunk is stored as a DocumentChunk record in the database with its positional index and source metadata.

Embeddings are generated for all chunks in a single batch call to the sentence-transformers model, which produces a 384-dimensional vector for each chunk. These vectors are written to the SQLite embeddings table as JSON arrays and simultaneously added to the FAISS index. The FAISS index is then serialised to disk so that the index state survives application restarts. Once all chunks and embeddings have been persisted, the document status is updated to indexed and the user is notified.




Figure 3.2

3.3 Retrieval and Generation Pipeline

The retrieval and generation pipeline is the core functional component of ClauseIQ. It transforms a natural-language query into a context-grounded response by combining semantic search with local language model inference.

When a query is submitted, the system first passes it through the guardrails module, which checks for prompt injection patterns and jailbreak attempts using a pattern-matching approach. If the query is classified as unsafe, it is rejected immediately and a refusal response is returned without any retrieval or generation taking place.

If the query is accepted, the system performs intent classification to determine whether it is document-specific, a general legal question, or outside the scope of the application. This classification is performed through keyword analysis rather than a separate model call, which keeps latency low. Document-specific queries are those that refer explicitly

to an uploaded file or use language such as summarise, analyse, or what does this say. General legal queries reference known Indian legal concepts, statutes, or procedures. Out-of-scope queries are those that fall outside the legal domain entirely and are returned with an appropriate boundary message.

For accepted queries, the system generates a query embedding using the same sentence-transformers model employed during ingestion, ensuring that query vectors and document vectors occupy the same semantic space. The query vector is then used to perform a k-nearest-neighbour search against the FAISS index, which returns the top five chunks ranked by cosine similarity. These chunks, along with their source metadata, are passed to the generation module.

The generation module constructs a structured prompt that includes the retrieved passages as numbered context blocks and the user's original query. For document-specific queries, the system prompt instructs the language model to answer exclusively from the provided context and to cite the relevant passage. For general legal queries, the model is permitted to supplement the retrieved context with its trained knowledge of Indian law. The completed prompt is submitted to the locally hosted language model through the Ollama API, and the response is returned to the user along with the source passages that were used.

Citation extraction is performed on the generated response by parsing references to the numbered context blocks. Any cited passage is extracted, truncated to a readable length, and returned to the front end as a separate citations field, where it is rendered beneath the main response to support traceability.

3.4 Database and Storage Design

The storage layer is designed to preserve users, documents, text chunks, embeddings, conversation sessions, and messages in a form that supports efficient retrieval and consistent integrity. A relational schema is appropriate for these entities because they carry well-defined identifiers and predictable relationships to one another.

The User table is designed to store the username, hashed password, role, account status, failed login count, lockout timestamp, and creation date. The Document table is designed to store the original filename, storage path, SHA-256 hash, MIME type, file size, upload timestamp, indexing status, and a reference to the owning user. The

DocumentChunk table is designed to store the extracted text, its positional index within the parent document, a JSON metadata object containing source information, and a reference to the parent document. The Embedding table is designed to store the chunk identifier, the embedding vector as a serialised array, and the name of the model used to produce it. The ChatSession table is designed to store the session name, creation timestamp, and a reference to the owning user. The ChatMessage table is designed to store the speaker role, message content, timestamp, and a reference to the parent session.

In addition to the relational database, the system is designed to maintain a FAISS index as a binary file on disk. This file is to be loaded into memory at startup and kept synchronised with new embeddings as documents are indexed. If the file is absent at startup, the system is designed to rebuild the index from embeddings already stored in the database, ensuring that retrieval capability remains consistent with persisted data.

Uploaded files are to be stored in a dedicated directory outside the application source tree, named using a combination of their SHA-256 hash and original filename to ensure uniqueness and prevent any two documents from overwriting each other.

3.5 Security and Privacy Design

Security is a central design consideration in ClauseIQ because legal queries may contain sensitive personal or professional information, and users must be able to trust that their data is handled responsibly at every stage.

Password security is to be enforced through one-way hashing at registration. No component of the system is designed to read, display, or log a plain-text password at any point. Session identifiers are to be stored server-side rather than embedded in URLs or exposed to client-side scripts, reducing the risk of session hijacking. Account lockout after repeated failed attempts is designed to limit exposure to automated attacks while still allowing legitimate users to regain access after the lockout period expires.

Data isolation between users is to be enforced at the service layer. Every query that retrieves documents, chunks, sessions, or messages is designed to be filtered by the authenticated user's identifier, so that a request from one user cannot return data belonging to another regardless of how the request is constructed.

The guardrails module is designed to screen every incoming query for patterns associated with prompt injection and jailbreak attempts. Queries that match these

patterns are to be rejected before any retrieval or generation takes place, preventing adversarial manipulation of the pipeline.

Privacy is addressed at the architectural level through the decision to run all inference locally via Ollama. Under this design, no query, document, or generated response is transmitted to any external server. The security of user data is therefore not dependent on the policies or availability of a third-party provider, which is particularly important for users who upload documents containing confidential information.

File upload handling is designed to include format validation based on MIME type and extension, a file size limit enforced before saving, and SHA-256 based deduplication to prevent redundant entries from accumulating in the index.

3.6 User Interface Design

The user interface is designed to minimise friction and to keep the interaction flow legible at every stage. The primary consideration is that users arrive with a legal question or a document to understand, and the interface is designed to present only what is relevant to the current task without exposing the internal complexity of the system.

The main chat view is designed as a two-column layout. The larger left column is intended to contain the conversation thread and the query input field, while the narrower right column is intended to display the list of documents the user has uploaded along with their current indexing status. This layout is designed to keep document context visible without interrupting the conversation flow, which is particularly useful when a user is asking a sequence of questions about the same document.

Conversation messages are designed to be rendered in visually distinct bubbles that differentiate user input from assistant responses. Where a response includes citations, the source passages are to be displayed beneath the main answer in a separate section, allowing the user to verify the grounding of the response without the citation text disrupting the readability of the primary content.

Document upload is designed to be accessible from the chat input area through an attachment control. Once a file is selected, the upload is to begin immediately and the document card in the right panel is designed to update in real time to reflect the current status, so that the user always knows whether their document is ready for querying.

Previous chat sessions are designed to be accessible through a sidebar that presents them in reverse chronological order, identified by their auto-generated name and creation date. Selecting a session is designed to restore the full conversation thread, allowing the user to continue from any prior point in their history.

The interface is to be implemented using Tailwind CSS utility classes applied within the HTML templates, keeping visual styling close to the markup it governs and avoiding the need for a separate stylesheet build process. The layout is designed to be responsive across standard laptop and desktop screen dimensions.

3.7 Data Flow Diagram (DFD)


The Data Flow Diagram of ClauseIQ illustrates how user inputs are processed and transformed into meaningful legal insights through a structured system architecture. The design clearly separates external entities, internal processing, and data storage components.

3.7.1 DFD Level 0

At this level, the ClauseIQ system is represented as a single central process interacting with two external entities: the user and the admin.

The user interacts with the system by submitting queries and uploading documents. In response, the system provides answers or legal insights. The admin interacts with the system for managing datasets, monitoring performance, and reviewing system reports.

The system is connected to two primary data stores. The vector database is used to store embeddings and perform similarity-based retrieval, while the metadata store maintains document information, logs, and historical records. This level provides a high-level overview without exposing internal processing details.



DFD Level 0


3.7.2 DFD Level 1

At this level, the internal workflow of ClauseIQ is expanded to show how data moves through the system.

The process begins when the system receives either a user query or an uploaded document. Uploaded documents are processed by extracting text and dividing it into smaller chunks suitable for analysis. These chunks are then converted into embeddings, which are stored in the vector database along with associated metadata.

When a query is submitted, the system retrieves relevant context by performing similarity search on the stored embeddings. The retrieved information, along with the original query, is passed to an external language model service to generate a response.

The generated answer is then formatted and returned to the user as a legal insight. Throughout this process, metadata such as document details, logs, and history are stored and managed in the metadata store.



DFD Level 1


3.8 System Flowchart

The system flowchart represents the step-by-step execution of ClauseIQ from input to output, including decision points and processing stages. The flow begins when a user either submits a query or uploads a document. If a document is uploaded, the system validates the file and proceeds to extract its content. In cases where direct extraction is not possible, an alternative extraction method is applied. The extracted text is then divided into smaller segments, and embeddings are generated and stored.

If the user submits a query, the system validates the input and retrieves the most relevant document segments using similarity search. These segments are combined with the query and passed to the language model to generate an answer.

Finally, the response is formatted and returned to the user along with relevant references. Error handling is incorporated throughout the process to ensure stability and a smooth user experience.

CLAUSEIQ – FLOW CHART



CLAUSEIQ Flow Chart

Chapter 4

Implementation and User Interface

4.1 Technology Stack and Environment

The implementation uses a stack selected to match the requirements of a small-to-medium scale intelligent web application while remaining reproducible on standard consumer hardware without cloud dependency.

The backend is written in Python 3.10 and built on the Flask micro-framework, which provides routing, request handling, session management, and templating through a lightweight and well-documented API. Flask-SQLAlchemy is used as the object-relational mapper, allowing the database schema to be defined as Python classes and queries to be expressed without raw SQL in the application logic. The database engine is SQLite, which requires no separate installation or server process and stores the entire relational dataset as a single file within the instance directory.

The retrieval layer relies on the sentence-transformers library, specifically the all-MiniLM-L6-v2 model, to generate 384-dimensional dense vector embeddings from document chunks and query strings. These embeddings are indexed and searched using FAISS, which performs inner-product similarity search in memory with sub-millisecond latency across tens of thousands of vectors. The FAISS index is serialised to disk after every update so that it survives application restarts without requiring documents to be re-indexed.

The generation layer communicates with a locally hosted language model through Ollama. Ollama manages model downloading, quantisation, and inference, exposing a REST API that the Flask backend calls with a structured prompt. The default chat model is configurable through the environment file, allowing the underlying model to be substituted without any change to the application code.

Document parsing uses pypdf for PDF files and python-docx for DOCX files. Optional OCR support is available through pytesseract and pdf2image for documents that contain scanned image content rather than machine-readable text. The frontend is rendered through Jinja2 templates served by Flask and styled using Tailwind CSS loaded from a CDN, with

interactive behaviour handled by vanilla JavaScript without any front-end framework dependency.

The development environment runs on a standard laptop without a dedicated GPU. All inference is performed on the CPU through Ollama's quantised model format, which makes the system deployable on hardware that is representative of the intended user base.

4.2 Backend Implementation

The backend is organised as a Flask application constructed through the application factory pattern, in which the `create_app` function initialises the database, registers blueprints, and configures extensions before returning the application instance. This pattern separates configuration from execution and makes the application straightforward to test and to start in different environments.

The routing layer is divided into four blueprints, each responsible for a distinct functional area. The `auth` blueprint handles registration, login, logout, session validation, and password management. The `upload` blueprint handles file receipt, validation, deduplication, and the dispatch of background indexing tasks. The `chat` blueprint receives query strings, invokes the retrieval and generation pipeline, and returns structured responses. The `pages` blueprint serves the HTML templates for each view.

Each blueprint contains route functions that perform only coordination - they validate inputs, call the appropriate service function, and format the response. Business logic is contained entirely within service modules that are imported by the routes. The `ingestion` service handles text extraction, chunking, embedding generation, and index updates. The `RAG` service handles query embedding, FAISS search, prompt construction, and Ollama invocation. The `guardrails` service handles query screening. This separation ensures that the backend remains readable and that individual service functions can be tested independently of the HTTP layer.

Background indexing is implemented using Python's `threading` module. When a document upload is accepted, a daemon thread is started that runs the ingestion pipeline against the saved file while the HTTP response is returned to the user immediately. The thread updates the document status in the database as it progresses, and the frontend polls the document status endpoint to reflect changes without requiring a page reload.

All protected routes check for an authenticated session before processing. If the session is absent or invalid, the request is redirected to the login page. Admin routes additionally

verify that the session role is set to admin, returning a 403 response to any authenticated user who does not hold that role.

4.3 Frontend Implementation

The frontend is implemented as a set of Jinja2 templates rendered server-side by Flask and enhanced with vanilla JavaScript for dynamic behaviour. No front-end framework is used, which keeps the dependency surface small and makes the templates straightforward to read and modify.

The base template defines the HTML document structure, loads Tailwind CSS from the CDN, and provides named blocks that child templates override for page-specific content and scripts. Each page template extends the base and fills these blocks with the markup relevant to that view. This inheritance structure prevents duplication of common layout elements across pages.

The chat template implements the primary user interface. It renders the two-column layout described in the design chapter, with the conversation thread and input field on the left and the document list panel on the right. JavaScript attached to the chat form intercepts the form submission event, reads the query value, sends it to the chat API endpoint as a JSON POST request, and appends the response to the conversation thread without reloading the page. User messages and assistant responses are rendered as distinct styled elements within the thread container. If the response includes a citations field, the citation passages are rendered beneath the main response text in a visually differentiated block.

The document panel is populated by a JavaScript function that calls the documents API endpoint and renders a card for each document returned. Each card displays the document name, file size, and current indexing status. A delete button on each card sends a DELETE request to the document endpoint and refreshes the panel on success. The panel is refreshed automatically when the page loads and after any upload or deletion event.

File upload is handled by a hidden file input element triggered by a paperclip button in the chat input area. When a file is selected, JavaScript constructs a FormData object and sends it to the upload endpoint. The document panel is refreshed after the response is received, and the user can observe the document status progressing from awaiting index to indexed as the background thread completes its work.

4.4 Upload and Indexing Workflow

The upload and indexing workflow transform a raw file submitted by the user into a set of searchable vector representations stored in both the relational database and the FAISS index. The workflow is designed so that each step either completes successfully or reports a specific failure, allowing the document status to reflect the outcome accurately at every stage.

When the upload endpoint receives a file, it first checks that the file extension and MIME type match the supported formats of PDF and DOCX. It then reads the file contents into memory to compute a SHA-256 hash, which is used to check whether an identical document has already been indexed. If a matching hash exists in the database, the upload is rejected and the user is informed of the duplication. If the document is new, the file is saved to the uploads directory under a path that incorporates the hash value, and a Document record is created in the database with the status set to awaiting index.

A background thread is then started to carry out the indexing process. The thread reads the saved file and passes it to the appropriate text extraction function based on its extension. For PDF files, `pypdf` iterates over the pages and extracts the text content of each. For DOCX files, `python-docx` iterates over the paragraph objects and joins their text.

Once text is available, it is divided into chunks. The chunking function splits the text at paragraph boundaries and accumulates paragraphs until the chunk reaches approximately eight hundred characters, at which point a new chunk is started. This boundary-aware splitting preserves the semantic coherence of paragraphs and avoids cutting a sentence mid-way. Each chunk is stored as a DocumentChunk record in the database with its positional index and a metadata object containing the source filename and document identifier.

Embeddings are generated for all chunks in a single call to the sentence-transformers model, which returns a list of vectors in the same order as the input chunks. Each vector is stored as an Embedding record in the database and simultaneously added to the in-memory FAISS index. After all embeddings have been added, the FAISS index is serialised to disk to ensure persistence. The document status is then updated to indexed, completing the workflow.

4.5 Conversation History and Personalisation

Conversation history provides continuity across sessions and supports the experience of working within a persistent application rather than a stateless query tool. Every interaction is stored against the authenticated user's account so that prior exchanges remain accessible after the browser is closed and reopened.

When a user submits their first query in a new context, a ChatSession record is created in the database. The session is assigned a name derived from the first thirty characters of the opening query, which provides a recognisable label in the session list without requiring the user to name conversations manually. All subsequent messages in that context are stored as ChatMessage records linked to the session, with a role field that distinguishes user messages from assistant responses and a timestamp that preserves the order of the exchange.

Previous sessions are retrieved and rendered in the sidebar of the chat interface in reverse chronological order. Each entry in the sidebar displays the session name and the date of its most recent message. Selecting an entry loads the full message history of that session into the conversation thread, allowing the user to review prior exchanges or continue the conversation from where it was left.

User-level data isolation is enforced throughout the history module. The query that retrieves sessions applies a filter on the user identifier stored in the session record, so that the sidebar never displays conversations belonging to a different account regardless of how the request is constructed. The same isolation applies to individual message retrieval.

4.6 Error Handling and Logging

Error handling is implemented at two levels: user-facing feedback delivered through the interface, and server-side logging recorded for debugging and maintenance purposes.

At the user-facing level, every API endpoint returns a structured JSON response that includes a success flag and, in the case of failure, a human-readable error message. The frontend JavaScript reads this message and renders it in the interface rather than displaying a raw HTTP error code. Invalid file types, oversized uploads, empty queries, duplicate documents, and authentication failures each produce a specific message that describes the problem and, where appropriate, indicates what the user should do differently.

Background operations such as document indexing update the document status field in the database to reflect failure states as well as successful completion. If text extraction yields no usable content, the status is set to a value that the frontend renders with an explanatory

label, informing the user that the document could not be indexed and suggesting that they verify the file is text-based rather than a scanned image.

At the server-side level, Python's built-in logging module is used to record significant events including document upload attempts, indexing start and completion, FAISS index updates, Ollama API calls, and authentication events such as failed login attempts and account lockouts. These records are written to a log file that is separate from the application source and can be reviewed without accessing the database directly. Log entries are structured to include a timestamp, a severity level, and a descriptive message, which makes them useful for tracing the sequence of events that preceded an unexpected outcome.

4.7 User Experience Refinements

Beyond the core query-answer loop, several interface refinements are designed to make the application feel complete and predictable for users who are not familiar with AI systems.

The Ollama health status is displayed as a small indicator in the application header. This indicator calls a lightweight health check endpoint when the page loads and at regular intervals, reporting whether the language model service is reachable. If Ollama is not running, the send button in the chat input is disabled and a visible message informs the user that the AI component is offline, preventing the frustration of submitting a query and waiting for a timeout response.

Document indexing progress is communicated through the document panel, where the status label updates from awaiting index to indexed as the background thread completes its work. The panel refreshes automatically at a short interval after an upload is received, so the user does not need to reload the page to see when a document becomes available for querying.

A legal disclaimer is displayed on the first visit to the chat interface, informing the user that ClauseIQ provides informational assistance based on retrieved legal text and does not constitute professional legal advice. The user acknowledges this notice before proceeding, which establishes the appropriate context for interpreting the system's responses.

The send button is disabled while a query is being processed and re-enabled when the response is returned, preventing the submission of duplicate requests during the response

latency period. The input field retains focus after each exchange so that the user can continue typing without manually clicking back into it.

Keyboard shortcuts are supported in the chat interface. The Enter key submits the current query, the Escape key clears the input field, and the focus is returned to the input automatically after a response is rendered. These shortcuts reduce the number of mouse interactions required during a continuous querying session.

4.8 Representative Screens and Observations

The screens presented in this section illustrate the system as it operates across its primary interaction flows. They are included to demonstrate that the design decisions described in Chapter 3 are reflected in the working interface and to provide a basis for the evaluation discussion in Chapter 5.

The landing page presents the application name, a brief description of its purpose, and a primary call-to-action that directs the user to the chat interface. The page is intentionally minimal, establishing the identity of the application without requiring the user to navigate through introductory content before reaching the functional interface.

The login and registration screens share a consistent single-column layout centred on the page. The registration form collects a username and password, with inline validation feedback displayed beneath each field if the submitted values fail the server-side checks. The login form provides the same layout with the addition of a link to the password recovery flow.

The chat interface occupies the full viewport and is divided into the two-column layout described in Chapter 3. A representative interaction in which a user submits a general legal query demonstrates the conversation thread rendering user and assistant turns in sequence, with the assistant response followed by a citation block that identifies the source passage used in generating the answer.

The document panel shows a representative state in which one document has been successfully indexed and a second is currently in the awaiting index state, illustrating the

asynchronous nature of the upload workflow and the status communication mechanism.




Figure 4.1: Landing Page




Figure 4.2: Legal Ai




Figure 4.3: Representative User Interface Layout

4.9 Project Execution Evidence

The following figure shows the geo-tagged image captured during the project development phase, indicating the team's presence and active participation.



Figure 4.4: Geo-tagged Image of Project Team

Chapter 5

Software Testing and Result Analysis

5.1 Test Planning

Testing was planned around the functional requirements established in Chapter 2. The test plan identifies the scenarios most likely to be encountered by a real user, as well as boundary conditions and error states that the system must handle without crashing or producing misleading output.

The areas covered by the test plan are user authentication, document upload and indexing, query submission and response generation, concurrent or rapid interaction, conversation history, session management, and system behaviour when a dependent service is unavailable. Each test case specifies a scenario description, the input or action taken, the expected outcome, and the actual observed result. Test cases are grouped into black-box tests, which examine the system from the user's perspective, and robustness tests, which examine behaviour under adverse or boundary conditions.

5.2 Black Box Test Cases

Black-box testing evaluates the system from the user's perspective without reference to internal implementation. Each test case describes an interaction that a user might reasonably attempt and records whether the system responded as expected.

The following table presents the full set of black-box test cases applied to the system.

ID	Test Scenario	Input / Action	Expected Result	Actual Result	Status
TC-01	Valid query submission	User enters a valid legal query and clicks submit	System processes query and returns relevant answer	Answer displayed correctly with references	Pass
TC-02	Empty query submission	Click submit without entering text	System should not process query and prompt user to enter text	Prompt displayed; no processing	Pass
TC-03	Irrelevant query	User submits unrelated query (e.g., cooking recipe)	System declines and shows out-of-scope message	Refusal message displayed	Pass
TC-04	Long query input	User enters very long query text	System handles input and returns response without crash	Response generated successfully	Pass
TC-05	Special characters in query	Query with symbols (@#\$\$%^&)	System sanitizes input and processes safely	No error; response generated	Pass
TC-06	Valid PDF upload	Upload supported PDF file	File accepted and indexed	Document appears and processed	Pass
TC-07	Unsupported file type	Upload .txt or .exe file	System rejects unsupported format	Error message displayed	Pass
TC-08	Large file upload	Upload file exceeding size limit	System rejects file with size error	File not uploaded	Pass
TC-09	Query on uploaded document	Ask question based on uploaded PDF	System retrieves answer from document	Correct context-based answer returned	Pass
TC-10	Multiple queries	User submits multiple queries sequentially	System handles each query independently	All responses correct	Pass
TC-11	Network interruption	Disconnect during query submission	System should handle failure gracefully	Error/timeout message shown	Pass
TC-12	Response accuracy	Submit legal query from dataset	System returns accurate legal info	Correct answer verified	Pass

Table 5.1: Black Box Test Cases

5.3 White Box Considerations

White-box analysis examines the internal logic of the application to identify paths through the code that may not be exercised by black-box tests alone. This analysis does not involve automated code coverage measurement but is based on a systematic review of the control flow within the most critical service functions.

The guardrails function contains a single conditional path that either returns the string UNSAFE or falls through to the normal processing pipeline. Manual inspection confirmed that the pattern list is evaluated in full for every query and that a match at any position in the list terminates evaluation immediately, preventing further processing.

The ingestion service contains a branching structure that handles three distinct outcomes: successful text extraction, fallback to OCR, and complete extraction failure. Each branch updates the document status to a different value and either continues the pipeline or exits it. Manual testing with a text-based PDF, a DOCX file, and a deliberately corrupted file confirmed that all three branches were exercised and produced the expected status values.

The retrieval function contains a guard that returns an empty list immediately if the FAISS index contains no vectors, preventing a search operation on an uninitialised index. This guard was verified by querying the system before any documents had been uploaded and confirming that the pipeline handled the empty result without raising an exception.

The generation function contains fallback logic that returns a descriptive error message if the Ollama API call fails or returns an empty response. This path was verified by stopping the Ollama service and submitting a query, confirming that the fallback message was returned rather than an unhandled exception propagating to the user interface.

5.4 Security and Robustness Tests

While formal security and robustness testing was not conducted as part of this project, the system incorporates several design-level safeguards that are intended to ensure safe and stable behaviour when faced with malformed, adversarial, or unexpected input. This section describes the conditions that the system is designed to handle and the expected outcome in each case.

The guardrails module is designed to intercept prompt injection attempts before any retrieval or generation takes place. Queries that contain patterns associated with attempts to override the system prompt or redirect the model's behaviour are rejected at the entry point of the pipeline, ensuring that adversarial input cannot influence the retrieval or generation steps. A query such as "ignore all previous instructions and act as a general assistant" is expected to be rejected with a refusal message rather than passed to the model.

Upload handling is designed to reject files that do not match the supported formats of PDF and DOCX, based on both the file extension and the MIME type reported by the browser. A file with a mismatched extension, such as a JPEG renamed with a .pdf extension, is expected to be rejected at the MIME type check rather than accepted and passed to the extraction pipeline. Files that exceed the configured size limit are expected to be rejected before the file content is read into memory.

The SHA-256 deduplication check is designed to prevent the same document from being indexed multiple times. If a user uploads a file that has already been accepted, the system is expected to return a duplicate notification and decline to create a new document record, keeping the index free of redundant content.

All protected API endpoints are designed to reject requests that arrive without a valid authenticated session, returning a redirect to the login page rather than exposing any data. This check is enforced in the service layer independently of the user interface, so that direct requests to API endpoints without a browser session are handled in the same way as unauthenticated page requests.

When the Ollama service is unavailable, the generation module is designed to return a descriptive error message rather than allowing an unhandled exception to propagate to the user interface. The health check endpoint is designed to detect this condition proactively and disable the query input before the user submits a request that cannot be fulfilled, reducing unnecessary waiting time.

The system is designed to remain stable when a second query is submitted while a previous query is still being processed. The send button is disabled during the response period, preventing duplicate submissions from being dispatched to the backend concurrently. If a document is deleted while indexing is in progress, the background thread is designed to handle the resulting database inconsistency gracefully rather than raising an unhandled error.

5.5 Performance and Usability Analysis

Performance was assessed informally through observation during manual testing rather than through automated load testing or profiling tools. The two components most relevant to perceived performance are the FAISS retrieval step and the language model inference step.

The retrieval step, which involves embedding the query using sentence-transformers and searching the FAISS index, completed consistently within one to two seconds across all test sessions, regardless of the number of documents indexed. This latency is attributable primarily to the embedding generation rather than the search itself, since FAISS operates in memory and the index sizes encountered during testing were well within the range where linear scan performance is negligible.

The generation step, which involves passing the structured prompt to the locally hosted language model through Ollama and waiting for the response, was the primary source of latency in the end-to-end pipeline. On the test hardware, which consists of a standard laptop without a dedicated GPU, responses were returned within fifteen to forty-five seconds depending on the complexity of the prompt and the length of the generated response. This latency is within the range that users found acceptable during informal observation, particularly given that the send button state and a visible status message communicate that processing is in progress.

Document indexing for files in the range of ten to thirty pages completed within thirty to ninety seconds in background threading, with no perceptible effect on the responsiveness of the chat interface during that period.

Usability was assessed by observing whether first-time interactions with the system produced the expected outcomes without requiring explanation. The health indicator in the header was found to be useful for communicating the availability of the language model before queries were attempted. The document status panel was effective in communicating indexing progress. The automatic naming of conversation sessions from the opening query was found to produce recognisable labels in all tested cases.

5.6 Results and Limitations

The primary limitation identified through testing concerns response quality rather than system correctness. When a query is phrased ambiguously or uses terminology that does not appear in the indexed corpus, the retrieval step may return passages that are topically adjacent but not directly relevant to the query. In these cases the language model generates a response that is coherent but less precise than it would be with a better-matched retrieval result. This limitation is inherent to the RAG architecture and is not specific to ClauseIQ; it reflects the general sensitivity of semantic retrieval to query phrasing.

A second limitation concerns the locally hosted language model. Smaller quantised models running on CPU-only hardware produce responses that are occasionally incomplete or less detailed than those produced by larger models on high-performance infrastructure. The system mitigates this by always providing retrieved context explicitly in the prompt, which reduces the degree to which the model must rely on its parametric knowledge, but it cannot fully compensate for the capacity difference between local and cloud-hosted models.

A third limitation concerns scanned documents. When an uploaded PDF consists of image-based pages rather than machine-readable text, the OCR fallback may produce imperfect extraction, particularly for documents with poor scan quality, non-standard fonts, or complex layouts. In such cases the indexed content may be incomplete, and the retrieval step may return less relevant passages than it would for a cleanly extracted document.

These limitations are acknowledged as boundaries of the current implementation rather than defects in the design

Chapter 6

Conclusion and Future Scope

6.1 Conclusion

This project demonstrates that retrieval-augmented generation can be applied meaningfully to a legal assistance context within the scope of an MCA final year project. ClauseIQ combines a structured document store, semantic vector search, a locally hosted language model, user authentication, and a conversational web interface into a coherent system that makes legal material more accessible to users without formal legal training.

The objectives defined in Chapter 1 have been addressed across the successive chapters of this report. The retrieval pipeline, built on FAISS and the sentence-transformers library, provides semantic search over indexed legal text and user-uploaded documents, producing more relevant results than keyword-based approaches for the query types tested. The generation pipeline, connected to a locally hosted language model through Ollama, produces responses that are grounded in retrieved passages rather than generated from model memory alone, which is an important property in a domain where accuracy and traceability matter. The document management workflow handles upload, extraction, chunking, embedding, and indexing through a background process that does not interrupt the conversational interface. Conversation history provides continuity across sessions, and the interface presents the system in a form that requires no technical knowledge to operate.

The system is not intended to replace professional legal counsel, and it does not attempt to do so. Its purpose is to reduce the difficulty of the first step that a non-specialist user faces when approaching a legal question - finding the relevant provision, understanding its language, and identifying what it requires. Within that defined scope, the project achieves its stated goals and produces a system that is deployable, testable, and extensible.

6.2 Future Scope

Several directions for future development have been identified through the course of the project. Each addresses a limitation noted in Chapter 5 or extends the system's capability in a direction that would increase its practical value.

The most immediate improvement concerns the chunking strategy. The current implementation divides text at paragraph boundaries and accumulates content up to a character limit. A more sophisticated approach would use sentence-aware splitting with controlled overlap, ensuring that no semantically complete statement is divided across two chunks. This change would improve retrieval precision for queries that map to short, specific provisions rather than to broad thematic sections.

A second direction concerns citation linking. The current system extracts source passages and presents them alongside the generated response, but it does not link these passages to their position within the original document or to an authoritative external source. A future version could associate each retrieved passage with a section number and a hyperlink to the official government publication of the relevant statute, which would allow users to verify the cited provision directly rather than relying solely on the indexed text.

A third direction involves the addition of a plain-language simplification feature. Users without legal training frequently encounter responses that, while accurate, still contain terminology that requires explanation. A secondary prompt pass that rewrites the generated response in simplified language, triggered by a toggle in the interface, would extend the system's accessibility without reducing the precision of the primary response.

A fourth direction concerns the seeding of the legal corpus from the structured data files already present in the application's data directory. The repository includes JSON files covering the IPC, CrPC, CPC, HMA, IEA, and BNS in structured form. Indexing these files at application startup as a system knowledge base, distinct from user-uploaded documents, would allow general legal queries to be answered from verified statutory text rather than from the model's parametric knowledge alone.

A fifth direction involves the replacement of the local SQLite database with PostgreSQL and the pgvector extension for production deployments. This change would allow embedding vectors to be stored and queried as native vector types rather than as JSON arrays, enabling database-level cosine similarity search and removing the dependency on a separate FAISS index file. It would also support concurrent multi-user access more robustly than the file-based SQLite engine.

Additional directions for future work include the introduction of feedback collection from users to support iterative improvement of retrieval ranking, the implementation of document comparison functionality to identify changes

between two versions of the same agreement, multilingual support to handle queries and documents in Hindi alongside English, and a more formal evaluation framework that measures retrieval precision and response accuracy against a defined test corpus rather than relying solely on manual observation.

6.3 Final Remarks

The development of ClauseIQ illustrates how applied software engineering can connect contemporary artificial intelligence techniques with a genuine informational need. The project traverses the complete lifecycle of a software product - from requirement analysis and stakeholder identification through design, implementation, testing, and reflection - and produces a system that is functional, documented, and suitable for further development.

The technical choices made throughout the project reflect a deliberate preference for transparency, reproducibility, and local operation. Every component of the stack is open-source, every processing step is traceable to a specific module, and no part of the pipeline requires access to a paid external service. These properties make the system appropriate not only as an academic submission but as a foundation for practical deployment in contexts where data privacy and infrastructure independence are important considerations.

The project also demonstrates that a student-developed system can operate at the intersection of natural language processing, information retrieval, web application engineering, and domain-specific knowledge without requiring specialised infrastructure or expert resources. The combination of these disciplines in a single coherent product is itself the primary contribution of the work, and the report captures both the technical design and the reasoning behind each decision.

Users of ClauseIQ are advised not to rely solely on the responses generated by the system when making decisions that carry legal consequences. The system retrieves passages from an indexed corpus and presents them through a language model that, while guided by the retrieved context, may produce responses that are incomplete, imprecisely worded, or insufficiently specific to the user's actual circumstances. Legal matters frequently involve procedural requirements, jurisdictional variations, and factual nuances that no retrieval-based system can fully account for without professional assessment. ClauseIQ is designed as a starting point for legal exploration - a tool for locating relevant provisions and understanding their general meaning - and not as a substitute for advice from a qualified legal practitioner. Any decision involving rights, obligations, penalties,

or formal legal proceedings should be made in consultation with a licensed advocate or legal professional who can evaluate the specific facts of the situation.

ClauseIQ, in its current form, is a guided assistant for legal exploration rather than a decision-making tool. That distinction is intentional and should be preserved in any future development. The value of the system lies in making source material easier to find, easier to read, and easier to question - not in replacing the judgment that qualified professionals apply to it.

6.2.1 Opportunities for Future Enhancement

Enhancement Area	Possible Direction
Citation support	Show the exact legal section or source used in the answer.
Role-based access	Distinguish between user, reviewer, and administrator privileges.
Feedback loop	Collect user ratings to refine retrieval quality.
Analytics dashboard	Summarize usage patterns and common topics.
Document comparison	Compare two legal sources or uploaded files side by side.

Table 6.1: Future Enhancement

APPENDIX A - DATA DICTIONARY

The following table documents the principal fields stored across the application database. Each entry identifies the table in which the field appears, its data type, its purpose within the system, and any constraints or remarks relevant to its correct interpretation. The table is included to provide a complete record of the stored data structure in a form suitable for academic submission and future development reference.

Field Name	Table	Data Type	Description	Remarks
session_id	Session	String (255)	Unique session identifier	Primary Key
name	Session	String (255)	Auto-generated user name	Optional
role	Session	String (20)	user / assistant	Default: user

Table 6.2: User / Session

Field Name	Table	Data Type	Description	Remarks
document_id	Document	Integer	Unique document ID	Primary Key
file_name	Document	String (255)	Name of uploaded file	Stored
file_type	Document	String (50)	Type (PDF, DOCX)	Validated
file_size	Document	Integer	Size in bytes	Limit enforced
sha256_hash	Document	String (64)	Unique file hash	Deduplication
upload_time	Document	Timestamp	Upload timestamp	Auto-generated

status	Document	String (20)	uploaded / indexed / failed	System controlled
--------	----------	-------------	-----------------------------	-------------------

Table 6.3: Document Metadata

Field Name	Table	Data Type	Description	Remarks
chunk_id	DocumentChunk	Integer	Unique chunk ID	Primary Key
document_id	DocumentChunk	Integer (FK)	Reference to document	Foreign Key
chunk_index	DocumentChunk	Integer	Order of chunk	Starts from 0
chunk_text	DocumentChunk	Text	Extracted text content	Preprocessed

Table 6.4: Document Chunks

Field Name	Table	Data Type	Description	Remarks
embedding_id	Embedding	Integer	Unique embedding ID	Primary Key
chunk_id	Embedding	Integer (FK)	Linked chunk	Foreign Key
vector	Embedding	JSON	Vector representation	High dimension
model_name	Embedding	String (100)	Embedding model used	e.g. BERT
created_at	Embedding	Timestamp	Creation time	Auto

Table 6.5: Embeddings

Field Name	Table	Data Type	Description	Remarks
query_id	Query	Integer	Unique query ID	Primary Key
session_id	Query	String (255)	User session	FK
query_text	Query	Text	User input query	Required
created_at	Query	Timestamp	Time of query	Auto
status	Query	String (20)	processed / failed	System

Table 6.6: Query Processing

Field Name	Table	Data Type	Description	Remarks
retrieval_id	Retrieval	Integer	Unique retrieval ID	Primary Key
query_id	Retrieval	Integer (FK)	Query reference	Foreign Key
chunk_id	Retrieval	Integer (FK)	Matched chunk	Foreign Key
score	Retrieval	Float	Similarity score	FAISS
rank	Retrieval	Integer	Ranking order	Top-k

Table 6.7: Retrieval Results

Field Name	Table	Data Type	Description	Remarks
response_id	Response	Integer	Unique response ID	Primary Key
query_id	Response	Integer (FK)	Linked query	Foreign Key
answer_text	Response	Text	Generated answer	LLM output
source_refs	Response	Text	Referenced chunks/docs	Traceability
generated_at	Response	Timestamp	Response time	Auto

Table 6.8: Response Generation

Field Name	Table	Data Type	Description	Remarks
error_id	ErrorLog	Integer	Unique error ID	Primary Key
query_id	ErrorLog	Integer	Related query	Nullable
error_type	ErrorLog	String (50)	Validation / System	Categorized
message	ErrorLog	Text	Error message	Debugging

timestamp	ErrorLog	Timestamp	Time of error	Auto

Table 6.9: Error Handling

APPENDIX B - SYSTEM WORKFLOW SUMMARY

The following narrative describes the end-to-end workflow of the ClauseIQ system as experienced by a typical authenticated user. It is included to provide a concise operational summary that complements the detailed design descriptions in Chapters 3 and 4.

The user navigates to the application landing page and is directed to the login screen. If they do not yet have an account, they complete the registration form by providing a username and password. On first login, the user is prompted to confirm their password. Subsequent logins proceed directly to the chat interface upon successful credential verification.

Once authenticated, the user lands on the chat dashboard. They may either submit a natural-language query immediately or upload a document before querying. If a document is uploaded, the system validates the file format and size, computes a SHA-256 hash to check for duplication, saves the file to the uploads directory, and dispatches a background thread to carry out text extraction, chunking, embedding generation, and FAISS index registration. The document panel on the right side of the interface displays the current indexing status and updates automatically as the background process progresses.

When the user submits a query, the system first passes it through the guardrails module to screen for prompt injection patterns. If the query is accepted, it is classified as document-specific, general legal, or out-of-scope. For accepted queries, the query text is embedded using the sentence-transformers model and the resulting vector is used to search the FAISS index for the most semantically similar chunks. The top-ranked passages are assembled into a structured prompt alongside the original query and submitted to the locally hosted language model through Ollama. The generated response is returned to the chat thread along with any cited source passages.

Every query and response is stored as a ChatMessage record linked to the current session. The session is named automatically from the first query. On subsequent visits, the user can access all previous sessions from the sidebar and resume any conversation from the point at which it was left. Documents and conversation history are isolated per user account and are not accessible to other registered users.

APPENDIX D - GLOSSARY OF TERMS

The following glossary defines the technical and domain-specific terms used throughout this report. Terms are listed in alphabetical order.

Term	Definition
BNS	Bharatiya Nyaya Sanhita - the Indian criminal code enacted in 2023 to replace the Indian Penal Code.
Chunk	A small segment of text produced by dividing a larger document during the ingestion pipeline, used as the unit of retrieval.
Cosine similarity	A measure of the angular distance between two vectors in a high-dimensional space, used to quantify the semantic closeness of a query to a document chunk.
CPC	Code of Civil Procedure - the primary legislation governing civil court procedure in India.
CrPC	Code of Criminal Procedure - the legislation governing criminal court procedure in India, partially replaced by the BNSS in 2023.
Deduplication	The process of detecting and rejecting duplicate document uploads based on the SHA-256 hash of the file contents.
Embedding	A dense numerical vector representation of a piece of text, generated by a language model, that encodes semantic meaning in a form suitable for similarity comparison.
FAISS	Facebook AI Similarity Search - an open-source library for efficient similarity search and clustering of dense vectors, used in ClauseIQ to index and retrieve document embeddings.
Guardrails	The screening module that checks incoming queries for prompt injection patterns and rejects unsafe inputs before they enter the retrieval and generation pipeline.
Indexing	The process of transforming extracted document text into embeddings and registering those embeddings in the FAISS index so that the document becomes searchable.
IPC	Indian Penal Code - the former primary criminal code of India, now partially superseded by the BNS.

JWT	JSON Web Token - a compact, self-contained format for representing claims between two parties, used in ClauseIQ for session token management.
LLM	Large Language Model - a neural network trained on large quantities of text that can generate coherent natural-language responses to prompts.
Ollama	A locally deployable service that manages the downloading, quantisation, and inference of open-source language models through a REST API.
ORM	Object-Relational Mapper - a software layer that allows database tables to be manipulated as Python class instances rather than through raw SQL queries; implemented in ClauseIQ using Flask-SQLAlchemy.
Prompt	The structured input sent to the language model, consisting of a system instruction, the retrieved context passages, and the user's original query.
Prompt injection	An adversarial technique in which a user attempts to override or bypass the system prompt by embedding instructions within their query input.
RAG	Retrieval-Augmented Generation - an architecture in which a language model's response is grounded in passages retrieved from an external knowledge source rather than generated solely from the model's parametric memory.
Session	The authenticated state maintained for a logged-in user across multiple HTTP requests, stored server-side and referenced through a secure cookie.
SHA-256	Secure Hash Algorithm 256-bit - a cryptographic hash function used in ClauseIQ to generate a fixed-length fingerprint of each uploaded file for deduplication purposes.
SQLite	A lightweight, file-based relational database engine that requires no separate server process, used as the primary data store in ClauseIQ.
Vector search	A retrieval method that identifies the most semantically relevant documents by comparing the numerical vector of a query against the vectors of indexed passages using a similarity metric.

Table D.1: Glossary of Terms

REFERENCES

1. Bharatiya Nyaya Sanhita, 2023.
2. Dey, L., et al. (2024). Retrieval-Augmented Generation for Knowledge Intensive Tasks. arXiv preprint.
3. Flask Documentation. Pallets Projects.
4. Python Documentation. Python Software Foundation.
5. SQLite Documentation. SQLite Consortium.
6. Ollama Documentation. Ollama.
7. The Code of Civil Procedure, 1908.
8. The Code of Criminal Procedure, 1973.
9. The Indian Evidence Act, 1872.
10. GeeksforGeeks.org